

INTRODUCTION

Model Setting and Notations

- K : Number of arms.
- δ : Tolerance level of wrong identification
- $\nu = \{\nu_a\}_{a=1}^K$: Reward distribution of arm a
- μ_a : Mean reward of arm $a \in [K]$
 $\Delta_{i,j} = |\mu_i - \mu_j|$, WLOG, $\mu_1 \geq \mu_2 \geq \dots \geq \mu_K$
- μ_0 : Known threshold for comparison
- $\mathcal{S}^{\text{pos}} = \{\nu : \mu_1 > \mu_0\}$, $\mathcal{S}^{\text{neg}} = \{\nu : \mu_1 < \mu_0\}$, For $\Delta > 0$,
 $\mathcal{S}_{\Delta}^{\text{pos}} = \{\nu : \mu_1 - \mu_0 \geq \Delta\}$, $\mathcal{S}_{\Delta}^{\text{neg}} = \{\nu : \mu_0 - \mu_1 \geq \Delta\}$.
- A_t : Action in round t
- X_t : Observed reward in round t
- $H(t) = \{(A_s, X_s)\}_{s=1}^t$: History collected **up to time t** .

Known Information

- K, μ_0, δ

Unknown Information

- ν_a, μ_a , for all $a \in [K]$

Dynamics and Model Uncertainty

At each round $t = 1, 2, \dots$, the Decision Maker

- 1 Pull an arm $A_t \in [K]$, A_t is $\sigma(H(t-1))$ -measurable,
- 2 Receive the outcome $X_t \sim \nu_{A_t}$,
- 3 Update the history $H(t) = H(t-1) \cup \{(A_t, X_t)\}$

The agent stops at the end of time step τ ,

τ is a **stopping time** with respect to the filtration $\{\sigma(H(t))\}_{t=1}^{\infty}$

Upon stopping, the agent **outputs arm $\hat{a} \in [K] \cup \{\text{None}\}$**

Definition of PAC Requirement

$i^*(\nu)$:

- $i^*(\nu) = \{a : \mu_a \geq \mu_0\}$, for $\nu \in \mathcal{S}^{\text{pos}}$
- $i^*(\nu) = \{\text{None}\}$, for $\nu \in \mathcal{S}^{\text{neg}}$

δ -PAC:

- A pulling strategy is δ -PAC, if for any $\delta \in (0, 1)$,
 $\nu \in \mathcal{S}^{\text{pos}} \cup \mathcal{S}^{\text{neg}}$, it satisfies $\Pr_{\nu}(\tau < +\infty, \hat{a} \in i^*(\nu)) > 1 - \delta$.

(Δ, δ) -PAC:

- A pulling strategy is (Δ, δ) -PAC, if it is δ -PAC, and for any $\Delta, \delta > 0$, we have $\sup_{\nu \in \mathcal{S}_{\Delta}^{\text{pos}} \cup \mathcal{S}_{\Delta}^{\text{neg}}} \mathbb{E}_{\nu} \tau < +\infty$.

Objective: Minimize $\mathbb{E}\tau$

The agent aims to design a (Δ, δ) -PAC pulling strategy (π, τ, \hat{a}) that **minimizes the *sampling complexity* $\mathbb{E}[\tau]$** .

Lit Review

Algorithm	Bound	Opt in pos	Opt in Neg
S-TaS	$\lim_{\delta \rightarrow 0} \frac{\mathbb{E}\tau}{\log \frac{1}{\delta}} = \begin{cases} H & \text{pos} \\ H_1^{\text{neg}} & \text{neg} \end{cases}$	✓	✓
HDoC	$\mathbb{E}\tau \leq \begin{cases} O(H \log \frac{K}{\delta} + H_1 \log \log \frac{1}{\delta} + \frac{K}{\epsilon^2}) & \text{pos} \\ O(H_1^{\text{neg}} \log \frac{K}{\delta} + \frac{K}{\epsilon^2}) & \text{neg} \end{cases}$	×	×
APGAI	$\mathbb{E}\tau \leq \begin{cases} O(H_0(\log \frac{K}{\delta})) & \text{pos} \\ O(H_1^{\text{neg}}(\log \frac{K}{\delta})) & \text{neg} \end{cases}$	×	✓
SEE (*)	$\mathbb{E}\tau \leq \begin{cases} O(H \log \frac{1}{\delta}) + O(H_1^{\text{pos}} \log \frac{K}{\Delta_{0,1}}) & \text{pos} \\ O(H_1^{\text{neg}}(\log \frac{1}{\delta} + \log H_1^{\text{neg}})) & \text{neg} \end{cases}$	✓	✓
Lower Bound (*)	$\mathbb{E}\tau \geq \begin{cases} O(H \log \frac{1}{\delta} + \frac{1}{m} H_1^{\text{low}} - \frac{1}{\Delta_{i,m+1}^2}) & \text{pos} \\ \Omega(H_1^{\text{neg}} \log \frac{1}{\delta}) & \text{neg} \end{cases}$	NA	NA

(*) denotes the result is from this paper

Extra Comments

- S-TaS (Degenne & Koolen 2019): only achieves asymptotic optimality
- HDoC (Kano et.al 2018), APGAI (Jourdan et.al 2023) are **not (Δ, δ) -PAC**
 - Two-armed instance
 - Arm 1 Gaussian, $\mu_1 > \mu_0$, Arm 2 Constant, $\mu_2 = \mu_0$
 - $\text{UCB}_1 < \mu_0$ holds with non-zero prob, then HDoC and APGAI will never stop
- SEE is nearly optimal
 - $\nu \in \mathcal{S}^{\text{neg}}$, all the arms are below μ_0
 - Only one arm is above μ_0
- For SEE, If $\nu \in \mathcal{S}^{\text{pos}}$, coefficient of $\log \frac{1}{\delta}$ is independent of arm number K

Algorithm

Sequential Exploration Exploitation (Informal)

- 1: **procedure** SEE(Input: Action set $[K]$, threshold μ_0 , tolerance level δ , $C > 1$).
- 2: **Tune** $\{\delta_k, T_k^{\text{et}}, T_k^{\text{ee}}\}_{k=1}^{+\infty}$.
- 3: **for** Phase $k = 1, 2, \dots$ **do**
- 4: (Exploration) Run algorithm LUCB_G with tolerance level δ_k and previous exploration history.
Stops until one of the two conditions holds.
 - Total pulling times in all exploration phases is not greater than T_k^{ee} .
Take $\hat{a}_k = \text{Not Complete}$.
 - LUCB_G stops and output $\hat{a}_k \in [K] \cup \{\text{None}\}$.
- 5: **if** $\hat{a}_k \in [K]$ **then**
- 6: (Exploitation) Keep pulling arm \hat{a}_k with independent samples
Stops until one of the two conditions holds.
 - Pulling times of \hat{a}_k is not smaller than T_k^{et} .
 - LCB defined by δ is above μ_0 , output \hat{a}_k as a qualified arm
- 7: **else if** $\hat{a}_k = \text{None}$ and $\delta_k < \frac{\delta}{3}$ **then**
- 8: Output the instance is negative
- 9: **end if**
- 10: **end for**
- 11: **return** x
- 12: **end procedure**

Notation on Complexity

$$H_1^{\text{neg}} = \sum_{a=1}^K \frac{2}{\Delta_{0,a}^2}, H_1^{\text{low}} = \sum_{a: \mu_a < \mu_0} \frac{2}{\Delta_{1,a}^2},$$

$$H_1^{\text{pos}} = \sum_{a=1}^K \frac{2}{\max\{\Delta_{0,a}^2, \Delta_{1,a}^2\}}, H = \frac{2}{\Delta_{0,1}^2}$$

$$H_1 = \sum_{a=2}^K \frac{2}{\Delta_{1,a}^2}, H_0 = \sum_{a: \mu_a \geq \mu_0} \frac{2}{\Delta_{0,a}^2},$$

$$H_1^{\text{BAI}} = \frac{2}{\Delta_{0,1}^2} + \sum_{a=2}^K \frac{2}{\Delta_{1,a}^2}.$$

Key Ideas Behind the Algorithm

Property of LUCB_G

- LUCB_G is an adapted BAI alg
 - Pull arms with highest UCB
 - Return an arm whose LCB is greater than μ_0
- Take $\delta' \in (0, 1)$ as input for LUCB_G.
Conditioned on event holds with prob $1 - O(\delta')$, LUCB_G
 - return $\mu_a > \mu_0$ with pulling times $O(H_1^{\text{pos}}(\log \frac{1}{\delta'} + \log H_1^{\text{pos}}))$
 - return **None**, with pulling times $O(H_1^{\text{neg}}(\log \frac{1}{\delta'} + \log H_1^{\text{neg}}))$
- Take $\delta' \in (0, 1)$ as input for LUCB.
If the event doesn't hold, LUCB might never stop

Conduct Exploration

- Call LUCB_G for exploration,
decreasing $\{\delta_k\}_{k=1}^{+\infty}$ as tolerance level for each phase
- Since it might get stuck in a non-stopping loop,
set up maximum pulling tiems T_k^{ee}

Conduct Exploitation, if $\hat{a}_k \in [K]$

- **Keep pulling \hat{a}_k with new samples**
- Output \hat{a}_k if its LCB(δ) $> \mu_0$

Accept $\hat{a}_k = \text{None}$

- Only when $\delta_k = \Theta(\delta)$
- Inspired by Degenne & Koolen 2019, negative instance is similar to Best Arm Identification

Numeric Settings

Benchmark Algorithms

- Adapted Murphy Sampling (Kaufmann et.al 2018)
- Adapted Track and Stop (Garivier & Kaufmann 2016)
- HDoC, LUCB_G (Kano et.al 2018)
- lilHDoC (Tsai et.al 2024)

Instance Setting

- All the mean rewards are below μ_0
- Fraction of Qualified Arms (100%, 50%, 25%, Unique)
- Linear reward vector

Numeric Experiments

Numeric Experiments

